Blowin' in the Wind:
The impact of wind farms on house prices

Samuele Giambra[*]

Brown University

May 2018

Abstract

This paper estimates the impact of wind turbines presence at the ZIP code and census tract level on house prices. A naïve OLS specification shows no evidence of a relation between wind turbines and home prices. Using an instrumental variables strategy where the instrument is defined as the percentage of the area of the ZIP code or census tract with potential wind capacity in 2008 greater than 30%, interacted with a dummy set to one in the years after the first wind turbine was installed in the US, I find that the presence of a wind farm significantly reduces house prices. An event-study design brings additional evidence of this relation.

Keywords: wind turbines, home prices, disamenities
JEL: D62, Q42, R30

1 Introduction

The world energy demand has grown exponentially over the last decades. In 2015 the overall energy consumption was about 575 quadrillion British thermal units (Btu) and the most recent projections from the 2017 International Energy Outlook forecast a new pick of 736 quadrillion Btu in 2040 (IEO 2017). Panel A of figure 1 shows that world energy use rose by 56.9% from 1990 through 2015.[1] In this framework, wind energy represents a key option for a sustainable future for two main reasons. First, wind is a clean and renewable energy source that produces virtually no toxic pollution or global warming emissions (Energy.gov 2017). Second, wind power production has very low operating and maintenance costs compared to traditional sources of energy, with nearly all the cost being related to turbines' construction and installation (Cullen 2013). In accordance with this view, the US wind energy sector has been characterized by a fast expansion during the last fifteen years, growing from a cumulative power capacity of less than 2,500 MW in 1999 to about 85,000 MW at the end of 2017 (see panel B of figure 1).

Even if the public opinion is typically in favor of wind power (Firestone & Kempton 2007), the construction of wind farms has often to overcome the opposition of local communities. A major concern is that the visual and aural impact of wind farms might decrease house values in the surrounding areas.[2] For instance, Slattery et al. (2012) conduct surveys in several communities near wind farms in Texas and Iowa and find that one third of the respondents consider wind turbines an unattractive feature of the landscape. In particular, the lowest levels of support are expressed by people living closest to wind farms, a phenomenon known as Not-In-My-Backyard (Swofford & Slattery 2010). From an international perspective, Ladenburg & Dubgaard (2007) estimate for Denmark that households are willing to pay between 46 and 122 Euros per year for siting a future offshore wind farm further away from the coast. For Australia, Hall et al. (2013) bring qualitative evidence of a "social gap" between publicly stated support and individual local acceptance, partly explained by the trade-off between global environmental gains and local bearing of the costs of visual changes to a place or landscape. According to Toke (2005), these are among the reasons that

---

[1]For a direct comparison with the previously reported fact, one million Btu is roughly equivalent to 25.1996 tonne of oil equivalent (toe).

[2]Another concern is that the noise and vibrations generated by wind farms could directly harm individuals living close to the plant. Although the traditional medical literature finds no evidence of a negative relationship between wind turbines and health (Knopper & Ollson 2011), a recent study claims that psychological factors might play a role in reports of sleep disturbance by people living in the proximity of wind plants (Jalali et al. 2016).

explain why less than half of onshore wind farm applications in England and Wales are successful through the normal planning process.

The object of this paper is to analyze the impact of wind turbines presence on house prices. Using data on the location and building date of each turbine from the Federal Aviation Administration (FAA), and home prices from Zillow and the Federal Housing Finance Agency (FHFA), I am able to estimate changes in house prices for areas that are or are not allocated a wind farm. However, the presence of wind turbines in a given geographic area is clearly an endogenous outcome since infrastructure projects are usually targeted towards areas in specific economic conditions[3] or influenced by political variables. I thus rely on two empirical strategies to try to isolate the causal impact of wind turbines on house prices. First, I instrument for wind turbines presence using the percentage of the area of a ZIP code or census tract with potential wind capacity in 2008 greater than 30%, interacted with a dummy set to one in the years after the first wind turbine was installed in the US. Gross potential wind capacity factors are estimated for the US by Lopez et al. (2012) taking into account the technology available in 2008 and local wind resources. The 30% threshold is often used by engineers to determine the profitability of a site (Kalmikov & Dykes 2010). The hypothesis underlying this instrument is that locations having a larger area above the threshold are more likely to be allocated a wind farm. Second, I perform an event study analysis to investigate how house prices change in the years before and after a wind farm is constructed. The event is defined as the time in which an aggregate of at least ten wind turbines becomes operative in a certain location.

The results from my preferred instrumental variable specification show that each additional wind turbine installed in a given year decreases home prices by about 0.25%. This is a remarkable negative effect considering that, on average, 17 new turbines are installed in a ZIP code/year, conditional on non-zero construction.[4] On the other hand, the results of the event study are more noisy due to the intrinsic difficulty to identify an exact point in time in which a wind plant starts exerting a negative impact on house prices. Consequently, this approach appears to be less informative about the causal impact of wind turbines on home prices.

This paper mainly contributes to a recent literature that analyzes the detrimental effects of

---

[3]Usually wind farms are built in rural areas and have the indirect aim of boosting the local economy (DOE/GO 2004).

[4]The standard deviation is about 27 new turbines in the average ZIP code/year.

wind turbines on house transaction prices. Sims et al. (2008) are among the first to study this topic, using a sample of 201 house transactions within a mile of a wind farm in Cornwall, UK. The authors do not find any significant evidence of a causal effect of wind turbines on house prices, but this conclusion might be driven by the limitations of the data used. More recently, Hoen et al. (2009) analyze the impact of 24 wind farms in the US using data on 7,500 sales of houses within 10 miles. In an improvement of this study, Hoen et al. (2013) use a differences-in-differences (DD) approach on more than 50,000 house transactions between 1996 and 2011 from areas within 10 miles of 67 wind farms in the US. Both of the aforementioned articles, however, are unable to uncover significant evidence of the impact of wind turbines on house prices. Focusing on Rhode Island, Lang et al. (2014) use data from 3,254 transactions of houses within a mile of a wind turbine, finding no significant effect. A fair amount of other studies employing DD techniques have instead reached conclusions in favor of the hypothesis of a negative impact on house value. For the Netherlands, Dröes & Koster (2014) use 2.2 million transaction prices and find that the construction of a wind turbine lowers the price of houses within 2km by 1.4 percentage points. Similarly, Sunak & Madlener (2016) find evidence of a negative impact of about 9-14% in two German cities. Gibbons (2015) develop a slightly more articulated digital elevation model to create a 200m grid that shows if a turbine is visible or not from each location. The author uses then a DD comparing price changes in locations where a new turbine is built and is visible with different comparison groups, such as price changes in places where the new turbine is hidden by the terrain or became visible in the past. The main conclusion of the study is that visible wind turbines decrease house prices by 5-6% within 2km, and by less than 1% within 14km.

In this paper, I contribute to this literature in two different ways. First, I enlarge the scope of the analysis using data on wind turbines in each state of the US through the end of 2017. Second, in the preferred specification, I address the endogeneity issue of the allocation of a wind turbine to a given geographic area using an instrumental variable approach similar to the one exploited by Dinkelman (2011), who instruments electrification projects in South Africa using land gradient. To my knowledge, this is the first attempt to use an IV approach to estimate the impact of wind turbines on house prices.

2 Data

2.1 Wind turbines

I use data on all wind energy projects in the US through December 2017, made available by the Federal Aviation Administration (FAA).[5] Wind turbines observations are divided into four categories:

- Determined with Built Date: records returned to the FAA with a Built Date; generally represent completed constructions.

- Determined without Built Date: records that are in some state of construction (and might be completed) but no Built Date has been returned to the FAA yet.

- Determined as Hazard: records that have been determined to be a hazard to air navigation and in general are not under construction.

- Not yet Determined: records in the early phase of the permitting process for which no determination of whether they are a hazard or not has been made.

In the main analysis, I use the sample of wind turbines that are Determined with Built Date. I drop observations that I am not able to match to any ZIP code or census tract (mainly because they are located in Alaska or offshore) or that are not classified as wind turbines.[6] I also remove observations below the first percentile of the height distribution to keep only turbines that have a significant impact on the local landscape. After imposing these restrictions, the sample of wind turbines for the analysis at the ZIP code level is comprised of 41,108 turbines, with a median height of 418 feet.[7;8]

I rely on the group of turbines Determined as Hazard for a falsification test. The categories Determined without Built Date and Not yet Determined are not used in this paper because of the higher difficulty in classifying each project as completed or not.

---

[5]The actual GIS files are publicly available for download from `https://www.fws.gov/southwest/es/Energy_Wind_FAA.html`.

[6]The data also include records for meteorological towers and other buildings unrelated to wind energy.

[7]Similarly, the sample of turbines for the analysis at the census tract level counts 43,575 turbines, with a median height of 416 feet.

[8]According to the NationalWindWatch (2006), two diffused turbines model are the 1.8-megawatt Vestas V90 from Denmark, with a total height of 410 feet, and the 2-megawatt Gamesa G87 from Spain, totaling 399 feet.

2.2 Potential wind capacity

Potential annual average gross capacity factors for the US are computed by Lopez et al. (2012) using data from the National Renewable Energy Laboratory (NREL). The authors use wind power at 80 meters height above surface to produce estimates of wind resources at a 200-meters horizontal spatial resolution. Next, they eliminate areas that cannot host wind plants (such as urban areas, federally protected lands, and onshore water features) and estimate average annual energy production assuming a power density of 5 MW/km^2 and 15% energy loss.[9] Finally, they divide the US map into cells with areas of approximately 400km^2 and compute the percentage of each cell having potential capacity greater than 30%.

Figure 2 shows that wind turbines (represented by red triangles) are mainly placed in locations with a higher proportion of their area with potential capacity of at least 30% (shaded in dark blue). High potential capacity factors tend to be concentrated in the so called "wind corridor", a region stretching from North Dakota and Montana southward to western Texas, and characterized by high average annual wind speeds.[10]

2.3 House prices

Data on house prices come from two distinct sources.

First, Zillow provides monthly prices on all single-family, condominium and co-operative homes for selected US ZIP codes.[11] I use two measures of house prices from Zillow: (i) median listing prices per square foot, and (ii) Zillow Home Value Index (ZHVI). Median listing prices are available from January 2010 through January 2018 for more than 10,000 distinct ZIP codes. On the other hand, ZHVI data span the period April 1996–January 2018 for more than 15,000 ZIP codes. The larger time and geographic coverage of the ZHVI makes it my preferred Zillow metric.

In order to take into account the changing composition of properties sold in different points in time, the Zillow Home Value Index is created from estimated sale prices (called Zestimates) on every home in Zillow's database instead of using prices from the subset of houses actually on

---

[9]These are common features of a typical wind technology available in 2008 (DOEEERE 2008).

[10]Average wind speed at 80 meters hight in the wind corridor ranges from 6 to 9 m/s (see map at `https://windexchange.energy.gov/maps-data/319`).

[11]Zillow datasets are publicly available for download from `http://www.zillow.com/research/data/`.

sale.[12] Appendix figure 1 shows the distribution of monthly log ZHVI, for prices below $1 million.

Second, I use the house price index (HPI) constructed by Bogin et al. (2016) (BDL henceforth) at the ZIP code and census tract level using data on 97 million home transactions from the Federal Housing Finance Agency (FHFA). The index is based on the repeat-sales methodology, which makes the assumption that the characteristics of a given house remain unchanged over the period of the analysis. Holding this true, the price index for a given location can be computed by regressing differences in sale prices of the *same* house on a vector of location×time dummies. Appendix section A provides more details about the repeat-sales methodology. A specific feature of BDL's index is the low level of geographic aggregation (ZIP code and census tract), obtained by aggregating sales occurring over a one-year horizon. In summary, BDL's index is available for approximately 18,000 ZIP codes and 55,000 census tracts from 1975 through 2017. Appendix figure 2 provides an overview of the ZIP codes for which this HPI exists in both 2000 and 2005.[13] After joining BDL house prices at the ZIP code level with wind turbines and wind power capacity data, the resulting sample consists of about 16,800 distinct ZIP codes, or more than 44,000 ZIP code-years. On the other hand, the number of wind turbines matched is just above 35%.

Figure 3 shows that, on average, the four price measures trace a similar pattern for the US housing market.

2.4 Additional variables

Home prices, turbines location and potential wind capacity are the central variables for the analysis that follows. I supplement this data with geographic and demographic controls, mainly at the ZIP code level.

First, I compute each ZIP code's distance from the coastline and elevation.[14] Second, average

---

[12]Zestimates are calculated daily for more than 100 million homes, using automated valuation models that are re-trained three times a week using updated data. Zestimates are then seasonally adjusted, and the ZHVI is defined as the median of all Zestimates in a given location (see `https://www.zillow.com/research/zhvi-methodology-6032/` for more details on ZHVI methodology and `https://www.zillow.com/zestimate/#acc` for an overview of Zestimates accuracy).

[13]Obviously, ZIP codes included in the sample are not randomly selected and they likely correspond to more populated geographic areas where home transactions occur more frequently. However, in most cases, they represent a relevant fraction of the complete set of each state's ZIP codes.

[14]The US coastline can be accessed from `https://earthworks.stanford.edu/catalog/stanford-xv279yj9196`, while the shapefile containing elevation data can be accessed through the online ArcGIS catalog. Distance from the coast and elevation are calculated from the ZIP code centroid.

income is computed as the ratio between adjusted gross income and number of returns in the 2010 tax year according to the IRS Statistics of Income series.[15] A location specific indicator of rurality is obtained from the Population Studies Center of the University of Michigan.[16] It is based on the Rural-Urban Continuum Codes computed by the United States Department of Agriculture (USDA), which classify each county into nine groups according to its level of urbanization. The mapping from counties to ZIP codes assumes that if a ZIP code falls into more than one county, the characteristic of the county with the largest share of the ZIP code will be used. The remaining demographic controls – total population, male ratio, and different ethnicity ratios – come from the 2010 Census.[17]

3   Sample characteristics

Wind energy projects are usually targeted towards less developed areas. Figure 4 shows that, on average, ZIP codes that are allocated a wind farm (defined as a total number of at least ten turbines) at some point in time have lower house prices than ZIP codes where a wind farm is never built. The difference between the two sub-samples is small (about $1,500) but remarkably stable over time.

Table 1 presents a summary of means and standard deviations of a battery of demographic and geographic variables. Column (1) displays descriptive statistics for the full sample, while column (2) and (3) distinguish between ZIP codes characterized by the presence of a wind farm or not. Areas with wind farms are on average more distant from the coastline, have a higher elevation, and are less populated. They have a lower average income, a higher fraction of white inhabitants, and lower fractions of blacks and asians. They also tend to have a higher fraction of male residents. The last row shows that locations that are not allocated a wind warm are characterized by a lower percentage of the area of the ZIP code with potential wind capacity in 2008 greater than 30%.

Column (4) of table 1 shows the results of t-tests on the equality of means in the two sub-samples, assuming unequal variances. ZIP codes with and without a wind farm are significantly

---

[15]Income data is publicly available from the IRS at
https://www.irs.gov/statistics/soi-tax-stats-individual-income-tax-statistics-zip-code-data-soi.
[16]Accessible at https://dsdr-kb.psc.isr.umich.edu/answer/1102.
[17]All variables are mapped to ZIP codes and made available by Jean Roth at http://www.nber.org/data/census-2010-zip-code-data.html.

different in terms of all variables considered. Since the following analysis relies on the percentage of the area of the ZIP code with potential wind capacity greater than 30% as excluded instrument, I further investigate if the covariates are balanced conditional on this variable. This can be seen as a partial test of the independence assumption of the instrumental variable design. Column (5) shows the coefficients estimated regressing each exogenous control on the percentage of the ZIP code area having potential wind capacity greater than 30%, while column (6) adds to the regression all the other covariates as controls, together with state fixed effects. The differences between the two sub-sample remain statistically significant, although the coefficients on the excluded instrument are relatively small. Overall however, the instrument does not seem to achieve a perfect balance of the geographic and demographic covariates considered. This informs the subsequent analysis.

4   Identification strategy

4.1   Instrumental variable approach

Annual average wind raster data with a resolution from 200-meter to 1000-meter cell sizes is made available by the .[18] For instance, appendix figure 1 shows the map assembled by NREL detailing wind level in Rhode Island. Wind speeds are recorded at a 50-meter height above the ground. Wind power is then grouped into classes that range from poor (class 1) to superb (class 7). High resolution land-based wind data for Alabama, Florida, Louisiana or Mississippi is unfortunately not available; these states will thus be excluded from the analysis.

4.2   Event study analysis

If the presence of a wind turbine in a given zip code was determined randomly, I could in principle be able to estimate its effect on house values using the following model

$$ln\left(p_{lt}\right) = \alpha_l + \beta turbines_{lt} + \mathbf{X}_{lt}\theta + \tau_t + \varepsilon_{lt} \tag{1}$$

---

[18]The data can be downloaded at `http://www.nrel.gov/gis/data_wind.html`.

where $ln(p_{lt})$ measures the average house value in location l (zip code or census tract), and time t (year or month), while $turbines_{lt}$ describes the number of wind turbines active in that location. The other parameters in the model – α_l, and τ_t – capture fixed effects for location and time, respectively. Finally, $\mathbf{X}_{lt}$ represents a vector of time varying location characteristics, such as average household income, total population, high education rate, and minority prevalence.

If we are willing to assume that the vector of location specific control variables remains constant over the time period relevant for this analysis – i.e. $\mathbf{X}_{lt} = \mathbf{X}_l$ – we can rewrite equation 1 in first differences

$$\Delta ln(p_{lt}) = \beta \Delta turbines_{lt} + \tau_t + \Delta \varepsilon_{lt}$$

However, the coefficients estimated running simple OLS on equation 1 are likely to be biased by the existence of omitted variables that are correlated with both the change in house values and wind turbine presence. For instance, if poorer geographic areas are more likely to be allocated a wind turbine, and the price of houses in these areas is lower, then this would generate a negative relationship between turbines and house prices. In the attempt to reduce this problem, I control for a vector of zip code level covariates, $\mathbf{X}_{zs0}$, measured in 2000, the baseline period of the analysis. Among the control variables I include distance from the coastline; elevation; percentage of urbanized area; total population; fraction of whites, blacks and asians; income per capita; and educational attainment. Unfortunately, even after controlling for these zip code characteristics it is still likely that the estimate of α_2 will be biased due to omitted or unobservable variables and confounding trends. To identify the causal impact of wind turbines on house values I therefore rely on an instrumental variable approach. I use annual average 50-meter height above the ground wind data to construct a variable that measures the percentage of the area of a given zip code having wind of power category equal or higher than three. Since the productivity of wind turbines depends, among other geographic characteristics, on the location's wind power, this instrument is likely to be a good predictor for wind turbine presence. The two equations related to the IV approach are

the following

$$\text{First stage:} \quad \Delta turbines_{lt} = \pi Z_{lt} + \tau_t + v_{lt} \tag{2}$$

$$\text{Second stage:} \quad \Delta ln\left(p_{lt}\right) = \beta \Delta \widehat{turbines}_{lt} + \tau_t + \Delta \varepsilon_{lt} \tag{3}$$

I will now discuss more in details why the instrument I construct is based on the cutoff of a power class equal to three. As mentioned above, wind data is reported by the NREL on the basis of seven wind power classes. In order to construct an instrumental variable that is relevant to the allocation decision of a wind turbine to a given location, I rely on the fact that, as noted in the Wind Energy Resource Atlas of the United States, "Areas designated class 3 or greater are suitable for most wind turbine applications, whereas class 2 areas are marginal. Class 1 areas are generally not suitable, although a few locations (e.g., exposed hilltops not shown on the maps) with adequate wind resource for wind turbine applications may exist in some class 1 areas."[19] Therefore, I choose to compute the percentage of the geographic area of each zip code that is characterized by average wind level of at least class three.[20] The results of the analysis that follows remain mainly unchanged if I choose slightly different specifications of the instrument, as for instance adding also areas with average wind of power class two.

5   Results

5.1   Threats to IV validity

I now present the main results of the IV approach outlined in the previous section. Table 2 reports the estimates obtained in the first stage. In all specifications standard errors are clustered at the date specific ($year \times month$) level. The dependent variable is an indicator that is equal to one whenever a wind turbine is present in a certain zip code and month. In column (1), the first row indicates that for a 10% increase in the area of the zip code having wind of at least category three, the probability of receiving a wind turbine increases by more than half percentage points. Column (2)

---

[19]The Wind Energy Resource Atlas of the United States can be consulted at `http://rredc.nrel.gov/wind/pubs/atlas/chp1.html`.

[20]Practically, this step is implemented in ArcGIS intersecting the map containing wind level information with that reporting the US zip code boundaries.

adds to this baseline model the vector of control variables listed before. The size of the coefficient remains basically unchanged, as well as the F-statistic which is highly above the rule-of-thumb threshold of 10. Moving to column (3), I include state and date specific fixed effects. In this case, the coefficient of the first row shows that a 10% increase in the area of a zip code with average wind level of power class of at least three increases by 0.6% the probability that a wind turbine will be assigned to that zip code. The coefficients of the covariates included in column (2) and (3) have generally the expected signs: for instance, zip codes characterized by high elevation have on average a higher probability to be allocated a wind turbine. On the other hand, wind turbines are less frequently found in richer and more populated areas. A somehow less intuitive result comes from the negative coefficient of the variable measuring the percentage of the area of the zip code being urbanized in year 2000. We would in fact expect that wind turbines are more frequently allocated to highly rural zip codes. A possible explanation to reconcile this finding is that wind turbines are not directly placed in urban areas but next to them, in order to represent an alternative source of energy in locations where energy demand is particularly high. Moreover, this could also be a mechanic consequence of the fact that the zip codes for which Zillow collects data on home selling prices are not a random selection of the universe of US zip codes.

Figure 5 is a graphic representation of the first stage just discussed. I first regress the instrument on the vector of covariates together with state fixed effects. Next, I take the residuals obtained from this regression and plot them against the indicator for turbine presence in the zip code. As a last step I fit a local quadratic regression on this data using the rule-of-thumb bandwidth suggested by Fan & Gijbels (1996). Consistent with the earlier estimates the probability to find a turbine in a given zip code is increasing in the residualized percentage area of the zip code with average wind of category three or higher. Consistently with the monotonicity assumption, also the slope is increasing in the residualized instrumental variable.

We turn next to the analysis of the results from the OLS regression of median home selling price per square foot on turbine presence, reported in the first three columns of Table 3. Standard errors are clustered at the date specific level in all specifications. The coefficients related to the complete model -- shown in column (3) -- suggest that the presence of a turbine in a given zip code decreases the house prices by $20, everything else kept constant. As intuitively expected, distance from the coastline and elevation have also a negative and significant effect on home selling prices,

while an opposite relation exists for instance in the case of income per capita.

Despite the OLS coefficient on wind turbine presence shows already some evidence of a negative effect on house prices, this relationship could be biased by the presence of omitted variables and endogeneity in the choice of the area where wind turbines are built. To be able to isolate an exogenous source of variation in wind turbines allocation to zip codes I instrument wind turbine presence with annual average wind level as discussed in the previous section. The coefficients estimated using the IV approach are reported in column (4) to (6). Once again, standard errors in parenthesis are clustered at the date specific level. In all three columns we confirm the hypothesis that wind turbine presence has a negative and statistically significant effect on home values. Focusing on column (6), where I include all controls together with state and date fixed effects, the negative impact on house prices is found to be as large as $122. It is important to notice that the 2SLS coefficients of wind turbine presence are way larger than the correspondent OLS estimates. This is likely due to the presence of measurement error in the price variable that is not completely eliminated by simply dropping outlier observations as discussed before.

6   Event study analysis

As an alternative approach to the instrumental variable estimation, I perform an event study of the effect of wind turbines on house values. In their study of displaced workers' earnings losses, Jacobson et al. (1993) develop a framework to analyze multiple events occurring at different timings. I adopt the notation presented in Sandler & Sandler (2012), who offer an overview of the recent literature of multiple event studies. Let y_{zst} denote the value of houses in zip code z, state s and time t, and call e^z the date of the event occurring in z. Then, I can estimate the following model

$$y_{zst} = \sum_{d=-D, d \neq -1}^{D} 1(t - e^z = d)\beta_d + \alpha_t + \gamma_s + \varepsilon_{zst}$$

where α_t denote time fixed effects and γ_s capture state fixed effects.

Figure 6 displays the coefficients β_d of the equation above for the median home selling prices per square foot collected by Zillow.

The time window considered is 3 years before and after a turbine is erected in a given zip code.

13

Observations in the plot are the coefficients estimated from the sample of zip codes where wind turbines are actually built. This implicit restriction imposed by the event study framework limits the power of this analysis since the dataset I use contains only 269 cases of zip codes where a wind turbine is located and information on house prices is available. Moreover, requiring to observe at least three years before and after the turbine is built limits this sample set even further. The consequence of these restrictions is that no clear pattern emerges from the event study plot shown in Figure 6. More precisely, the figure shows that home selling prices decrease on average the year before a wind turbine is built, but this drop is not different in magnitude than the one measured two years after the event making the interpretation of any of those pattern less credible. Moreover, the standard errors related to the coefficients β_d (not depicted in the figure) are extremely large and eliminate any statistical difference between the observations plotted.

7 Conclusions and future work

This paper investigates the relation between wind turbine presence at the zip code level and home values. It uses the percentage of the area of a given zip code having annual average wind level of category three or more to address problems related to the endogenous placement of wind turbines and other confounding trends. Results from the 2SLS estimation suggest that home selling prices in places that are allocated a wind turbine are on average $120 lower. On the other hand, the event study analysis does not convincingly show any particular trend before and after a wind turbine is constructed.

Despite the findings discussed in this paper clearly point towards the direction of an existing negative relation between wind turbines presence and home values, this evidence if far from being conclusive. However, Zillow will soon release a new version of its public available house price dataset that will cover the entire US at the census tract level. This will help reducing the issue of sample selection faced by this study and likely add more robustness to its results.

References

Bogin, A., Doerner, W. & Larson, W. (2016), 'Local house price dynamics: New indices and stylized facts', *Real Estate Economics* .

Cullen, J. (2013), 'Measuring the environmental benefits of wind-generated electricity', *American Economic Journal: Economic Policy* **5**(4), 107–33.

Dinkelman, T. (2011), 'The effects of rural electrification on employment: New evidence from south africa', *The American Economic Review* **101**(7), 3078–3108.

DOEEERE (2008), 20Increasing wind energy's contribution to u.s. electricity supply, Technical report, National Renewable Energy Laboratory.

DOE/GO (2004), Wind energy for rural economic development, Technical report, U.S. Department of Energy.

Dröes, M. & Koster, H. (2014), 'Renewable energy and negative externalities: the effect of wind turbines on house prices'.

Energy.gov (2017), 'Advantages and challenges of wind energy'.
URL: *https://www.energy.gov/eere/wind/advantages-and-challenges-wind-energy*

Fan, J. & Gijbels, I. (1996), *Local polynomial modelling and its applications: monographs on statistics and applied probability 66*, Vol. 66, CRC Press.

Firestone, J. & Kempton, W. (2007), 'Public opinion about large offshore wind power: underlying factors', *Energy policy* **35**(3), 1584–1598.

Gibbons, S. (2015), 'Gone with the wind: Valuing the visual impacts of wind turbines through house prices', *Journal of Environmental Economics and Management* **72**, 177–196.

Hall, N., Ashworth, P. & Devine-Wright, P. (2013), 'Societal acceptance of wind farms: Analysis of four common themes across australian case studies', *Energy Policy* **58**, 200–208.

Hoen, B., Brown, J. P., Jackson, T., Wiser, R., Thayer, M. & Cappers, P. (2013), 'A spatial hedonic analysis of the effects of wind energy facilities on surrounding property values in the united states'.

Hoen, B., Wiser, R., Cappers, P., Thayer, M. & Sethi, G. (2009), 'The impact of wind power projects on residential property values in the united states: A multi-site hedonic analysis'.

IEO (2017), International energy outlook 2017, Technical report, U.S. Energy Information Administration.

Jacobson, L. S., LaLonde, R. J. & Sullivan, D. G. (1993), 'Earnings losses of displaced workers', *The American economic review* pp. 685–709.

Jalali, L., Nezhad-Ahmadi, M.-R., Gohari, M., Bigelow, P. & McColl, S. (2016), 'The impact of psychological factors on self-reported sleep disturbance among people living in the vicinity of wind turbines', *Environmental research* **148**, 401–410.

Kalmikov, A. & Dykes, K. (2010), 'Wind power wind power fundamentals', University Lecture.

Knopper, L. D. & Ollson, C. A. (2011), 'Health effects and wind turbines: A review of the literature', *Environmental Health* **10**(1), 78.

Ladenburg, J. & Dubgaard, A. (2007), 'Willingness to pay for reduced visual disamenities from offshore wind farms in denmark', *Energy Policy* **35**(8), 4059–4071.

Lang, C., Opaluch, J. J. & Sfinarolakis, G. (2014), 'The windy city: Property value impacts of wind turbines in an urban setting', *Energy Economics* **44**, 413–421.

Lopez, A., Roberts, B., Heimiller, D., Blair, N. & Porro, G. (2012), Us renewable energy technical potentials: a gis-based analysis, Technical report, NREL.

NationalWindWatch (2006), 'How big is a wind turbine?'.
  URL: *https://www.wind-watch.org/faq-size.php*

Sandler, D. H. & Sandler, R. (2012), 'Multiple event studies in public finance: a simulation study with applications', *Working Paper* **5**, 1–27.

Sims, S., Dent, P. & Oskrochi, G. R. (2008), 'Modelling the impact of wind farms on house prices in the uk', *International Journal of Strategic Property Management* **12**(4), 251–269.

Slattery, M. C., Johnson, B. L., Swofford, J. A. & Pasqualetti, M. J. (2012), 'The predominance of economic development in the support for large-scale wind farms in the us great plains', *Renewable and Sustainable Energy Reviews* **16**(6), 3690–3701.

Sunak, Y. & Madlener, R. (2016), 'The impact of wind farm visibility on property values: A spatial difference-in-differences analysis', *Energy Economics* **55**, 79–91.

Swofford, J. & Slattery, M. (2010), 'Public attitudes of wind energy in texas: Local communities in close proximity to wind farms and their effect on decision-making', *Energy policy* **38**(5), 2508–2519.

Toke, D. (2005), 'Explaining wind power planning outcomes:: some findings from a study in england and wales', *Energy policy* **33**(12), 1527–1539.

A   Repeat-sales index methodology

This appendix investigates

Table 1: Covariates balance

Covariates	Means			Differences in means		
	Full sample	Wind farm	No wind farm	Columns (2)-(3)	By wind	
					No controls	Controls
	(1)	(2)	(3)	(4)	(5)	(6)
Average income	56.4481	47.5509	56.5857	-9.0347	-0.1097	-0.1116
	(34.3034)	(10.7874)	(34.5236)	(0.7299)	(0.0168)	(0.0168)
Total population	15055.1253	8941.6601	15149.6384	-6207.9783	-178.1270	-99.0650
	(14777.3287)	(10314.7706)	(14816.2206)	(658.7455)	(7.1227)	(6.4686)
Elevation	294.6454	503.6166	291.4148	212.2018	1.2884	-0.4558
	(383.4397)	(403.9627)	(382.2311)	(25.5721)	(0.1880)	(0.1023)
Distance to the coast	228.1892	379.1010	225.8561	153.2449	3.1656	0.3865
	(267.0270)	(320.8699)	(265.4482)	(20.2794)	(0.1288)	(0.0494)
Male ratio	0.4950	0.5019	0.4949	0.0069	0.0001	0.0000
	(0.0220)	(0.0233)	(0.0219)	(0.0015)	(0.0000)	(0.0000)
White ratio	0.8286	0.9082	0.8274	0.0809	0.0017	0.0001
	(0.1869)	(0.1021)	(0.1877)	(0.0066)	(0.0001)	(0.0000)
Black ratio	0.0833	0.0180	0.0843	-0.0663	-0.0011	0.0001
	(0.1520)	(0.0397)	(0.1528)	(0.0028)	(0.0001)	(0.0000)
Asian ratio	0.0242	0.0087	0.0244	-0.0157	-0.0003	0.0001
	(0.0514)	(0.0282)	(0.0516)	(0.0018)	(0.0000)	(0.0000)
Percentage ZIP code area with	3.5990	21.5823	3.3210	18.2613		
potential wind capacity > 30%	(15.8003)	(36.1620)	(15.1093)	(2.2766)		
Number of ZIP codes	16618	253	16365	16618	16618	16618

Notes: The sample is the set of ZIP codes having information on Bogin et al. (2016)'s house price index for at least one year and the full vector of control variables. See section 2 for each variable definition and source. Standard errors are reported in parenthesis. Column (1) reports covariates averages for the full sample of ZIP codes. Columns (2) and (3) divide the sample into two sub-groups, according to whether a cumulative number of at least ten turbines is ever observed in a given ZIP code or not. Column (4) tests for differences between column (2) and (3) using a t-test with unequal variances. Column (5) reports the coefficient obtained regressing each covariate on the percentage of the ZIP code area with potential wind capacity greater than 30%. Column (6) additionally controls for the remaining covariates.

Table 2: OLS regressions

	Panel A		Panel B		Panel C		Panel D	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
$\Delta ln(p)$								
Δturbines	0.00016	0.00016	0.00010	0.00011	0.00009	0.00005	-0.00003	-0.00002
	(0.00006)	(0.00006)	(0.00007)	(0.00007)	(0.00015)	(0.00016)	(0.00003)	(0.00003)
State fixed effects	No	Yes	No	Yes	No	Yes	No	Yes
Number of locations	46731	46731	16786	16786	9309	9309	14082	14082
Number of location-time	1147555	1147555	442092	442092	722194	722194	3309753	3309753

Notes: In panel A, a time period is a year and the unit of analysis is the census tract-year. In panel B, a time period is a year and the unit of analysis is the ZIP code-year. In panel C and D, a time period is a month and the unit of analysis is the ZIP code-month. House prices in panel A and B are estimated by Bogin et al. (2016). House prices in panel C are median listing prices per square foot from Zillow, while specifications in panel D use the Zillow Home Value Index. Each regression includes calendar time fixed effects. Standard errors in parenthesis are clustered by location.

Table 3: First stage regressions

	Panel A		Panel B		Panel C		Panel D	
$\triangle$turbines	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Percentage area with potential wind capacity > 30% x $\mathbb{1}$(year first turbine US)	0.0041 (0.0004)	0.0052 (0.0005)	0.0048 (0.0005)	0.0062 (0.0007)	0.0006 (0.0001)	0.0008 (0.0002)	0.0005 (0.0001)	0.0007 (0.0001)
State fixed effects	No	Yes	No	Yes	No	Yes	No	Yes
Number of locations	46731	46731	16786	16786	9309	9309	14082	14082
Number of location-time	1147555	1147555	442092	442092	722194	722194	3309753	3309753

Notes: In panel A, a time period is a year and the unit of analysis is the census tract-year. In panel B, a time period is a year and the unit of analysis is the ZIP code-year. In panel C and D, a time period is a month and the unit of analysis is the ZIP code-month. The excluded instrument is defined as the percentage of the area of a given location having potential wind capacity in 2008 greater than 30%, interacted with a dummy set to one in the years after the first wind turbine was installed in the US. Each regression includes calendar time fixed effects. Standard errors in parenthesis are clustered by location.

Table 4: IV regressions

$\triangle ln(p)$	Panel A		Panel B		Panel C		Panel D	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
$\triangle$turbines	-0.0220	-0.0109	-0.0091	-0.0025	-0.0059	-0.0149	-0.0197	-0.0025
	(0.0023)	(0.0013)	(0.0013)	(0.0008)	(0.0030)	(0.0041)	(0.0026)	(0.0009)
State fixed effects	No	Yes	No	Yes	No	Yes	No	Yes
Number of locations	46731	46731	16786	16786	9309	9309	14082	14082
Number of location-time	1147555	1147555	442092	442092	722194	722194	3309753	3309753

Notes: In panel A, a time period is a year and the unit of analysis is the ZIP code-year. In panel B, a time period is a year and the unit of analysis is the census tract-year. In panel C and D, a time period is a month and the unit of analysis is the ZIP code-month. House prices in panel A and B are estimated by Bogin et al. (2016). House prices in panel C are median listing prices per square foot from Zillow, while specifications in panel D use the Zillow Home Value Index. The excluded instrument is defined as the percentage of the area of a given location having potential wind capacity in 2008 greater than 30%, interacted with a dummy set to one in the years after the first wind turbine was installed in the US. Each regression includes calendar time fixed effects. Standard errors in parenthesis are clustered by location.

Table 5: Monotonicity test, Bogin et al. (2016) prices at ZIP code level

	△turbines	Coefficient (Standard error)	Number of locations	Number of location-time
(1)	Counties in metro areas of 1 million population or more	0.0009 (0.0004)	5654	176352
(2)	Counties in metro areas of 250,000 to 1 million population	0.0029 (0.0008)	3514	98398
(3)	Counties in metro areas of fewer than 250,000 population	0.0036 (0.0010)	2233	56479
(4)	Urban population of 20,000 or more, adjacent to a metro area	0.0043 (0.0012)	1429	32263
(5)	Urban population of 20,000 or more, not adjacent to a metro area	0.0058 (0.0034)	410	9493
(6)	Urban population of 2,500 to 19,999, adjacent to a metro area	0.0056 (0.0015)	1734	33185
(7)	Urban population of 2,500 to 19,999, not adjacent to a metro area	0.0107 (0.0017)	843	16598
(8)	Completely rural or less than 2,500 urban population, adjacent to a metro area	0.0042 (0.0028)	344	5493
(9)	Completely rural or less than 2,500 urban population, not adjacent to a metro area	0.0013 (0.0007)	281	4555

Notes: Each specification use a different sub-sample from the set of ZIP codes having information on house prices estimated by Bogin et al. (2016). Sub-samples are defined according to a measures of rurality at the county level ranging from 1 through 9 (see `https://dsdr-kb.psc.isr.umich.edu/answer/1102`). For each regression, the independent variable is the percentage of the area of a given ZIP code having potential wind capacity in 2008 greater than 30%, interacted with a dummy set to one in the years after the first wind turbine was installed in the US. All specifications include year fixed effects and standard errors are clustered at the ZIP code level.

Table 6: Falsification test

$\triangle ln(p)$	Panel A		Panel B		Panel C	
	(1)	(2)	(3)	(4)	(5)	(6)
Potential wind capacity	0.00008	0.00009	0.00000	0.00009	-0.00002	-0.00001
	(0.00001)	(0.00001)	(0.00001)	(0.00001)	(0.00000)	(0.00000)
State fixed effects	No	Yes	No	Yes	No	Yes
Number of locations	38202	38202	13322	13322	13958	13958
Number of location-time	364161	364161	157663	157663	791295	791295

Notes: The sample is restricted to the period before the first wind turbine was installed in the US. In panel A, a time period is a year and the unit of analysis is the census tract-year. In panel B, a time period is a year and the unit of analysis is the ZIP code-year. In panel C, a time period is a month and the unit of analysis is the ZIP code-month. House prices in panel A and B are estimated by Bogin et al. (2016). The specifications in panel C use the Zillow Home Value Index. The independent variable is defined as the percentage of the area of a given location having potential wind capacity in 2008 greater than 30%. Each regression includes calendar time fixed effects. Standard errors in parenthesis are clustered by location.

Figure 1: Energy trends

Panel A: World energy consumption


Panel B: US wind energy capacity

Notes: Data on world energy consumption comes from the 2016 Global Energy Statistical Year-book (`https://yearbook.enerdata.net/`). Data on US wind energy capacity comes from the Office of Energy Efficiency & Renewable Energy (`https://windexchange.energy.gov/maps-data/321`).

Figure 2: Wind turbines location and potential wind capacity in 2008
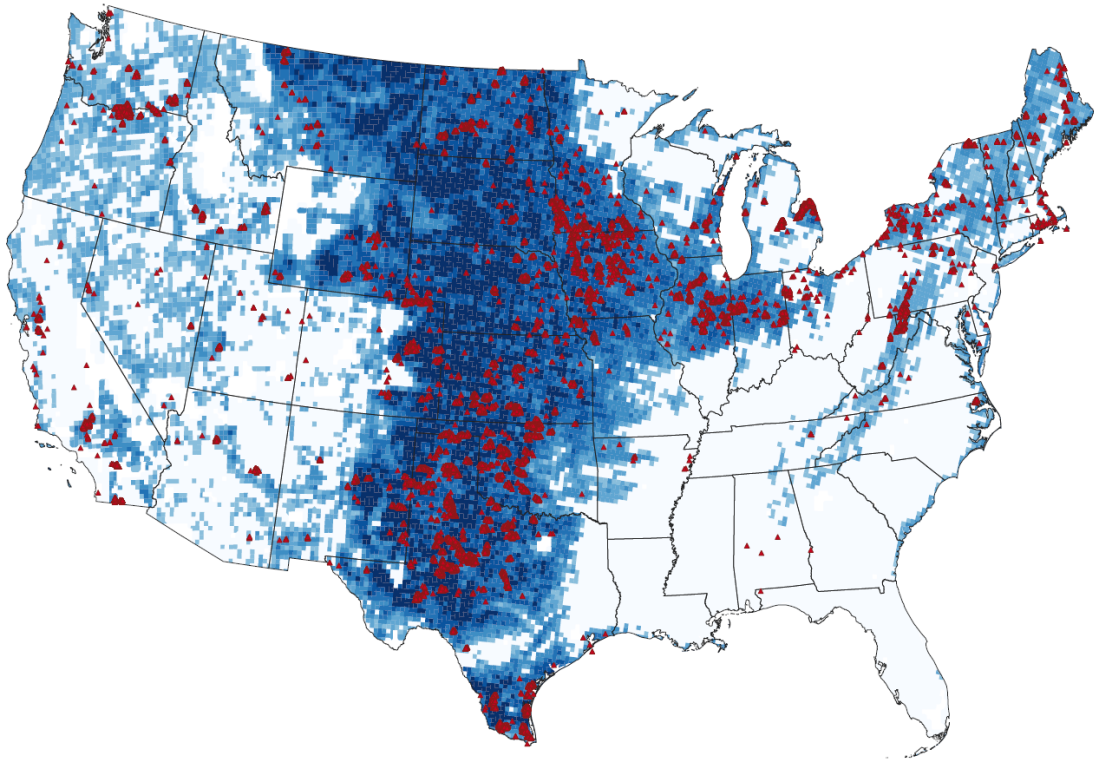
Notes: The figure shows data on potential wind capacity in 2008 estimated by Lopez et al. (2012) using data from the National Renewable Energy Laboratory (NREL). The US map (Alaska excluded) is divided into cells of approximately $400km^2$. Blue cells correspond to locations having a non-zero fraction of their area with potential wind capacity larger than 30%. The intensity of the color represents the proportion of the area with wind capacity above the threshold. Red triangles symbolize wind turbines installed by December 2017. Data on wind turbines comes from the Federal Aviation Administration (FAA).

Figure 3: Comparison of house prices measures
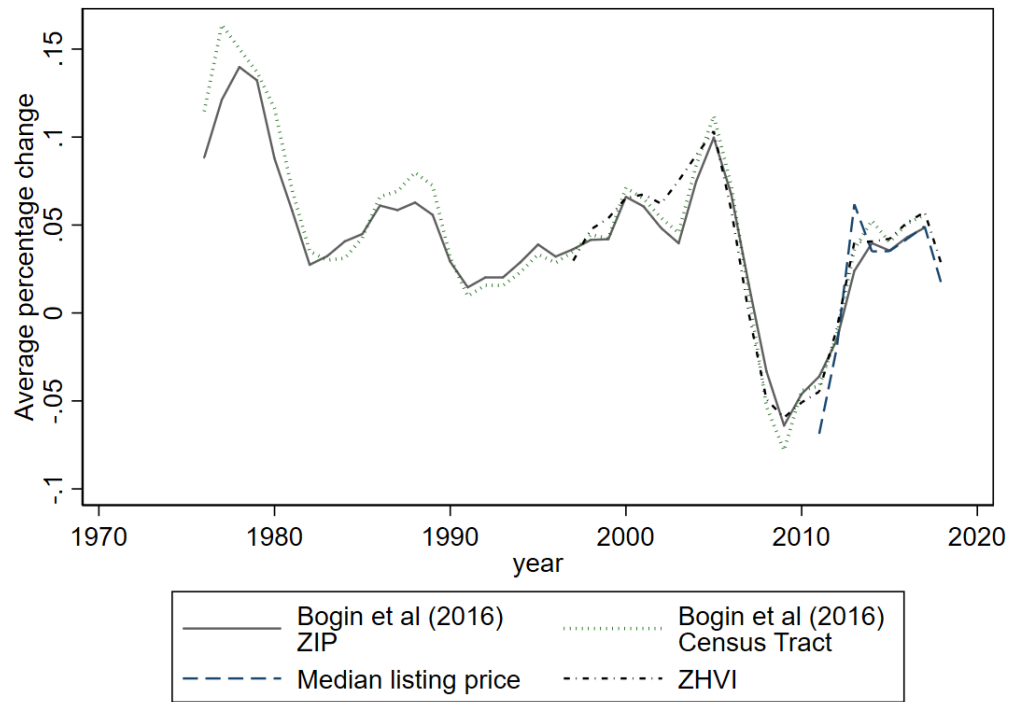
Notes: Data on home prices are either from Zillow or estimated by Bogin et al. (2016). Prices from Zillow are first averaged at the ZIP code-year level. Each series is then computed as the average percentage change between two consecutive years.

Figure 4: House prices time series

Notes: The figure reports data on average monthly log Zillow Home Value Index. The sample is divided into two sub-groups, according to whether a cumulative number of at least ten turbines is ever observed in a given ZIP code (solid line) or not (dashed line).

Figure 5: Graphic representation of instrumental variable approach

Panel A: First stage

Notes: Bandwidth is 8.029

Panel B: Reduced form

Notes: Bandwidth is 7.430
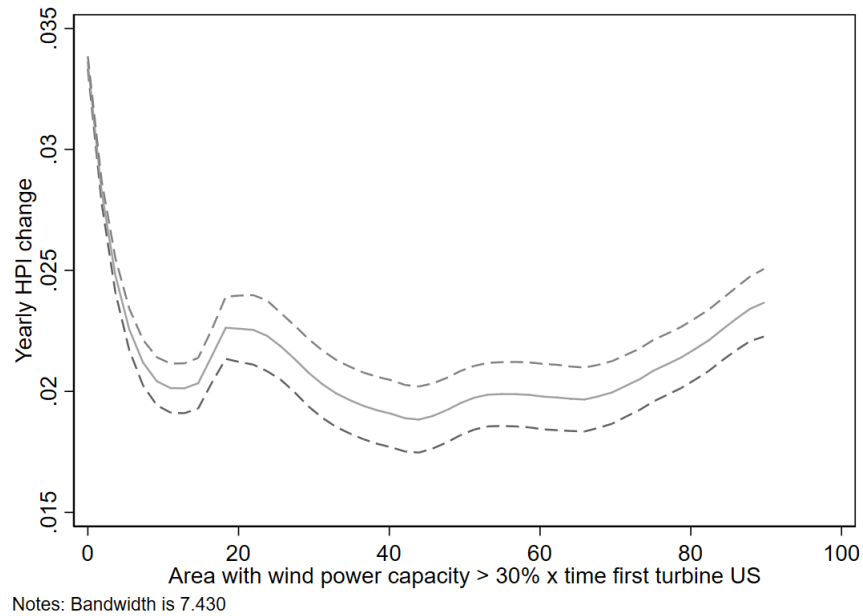
Notes:

Figure 6: Yearly house price

Notes:

Appendix Table 1: Monotonicity test, Bogin et al. (2016) prices at census tract level

|	$\triangle$turbines	Coefficient (Standard error)	Number of locations	Number of location-time
(1)	Counties in metro areas of 1 million population or more	0.0013 (0.0005)	21645	568137
(2)	Counties in metro areas of 250,000 to 1 million population	0.0035 (0.0012)	10759	270462
(3)	Counties in metro areas of fewer than 250,000 population	0.0044 (0.0011)	5661	135328
(4)	Urban population of 20,000 or more, adjacent to a metro area	0.0031 (0.0010)	2819	62017
(5)	Urban population of 20,000 or more, not adjacent to a metro area	0.0054 (0.0021)	984	21888
(6)	Urban population of 2,500 to 19,999, adjacent to a metro area	0.0039 (0.0009)	2647	49892
(7)	Urban population of 2,500 to 19,999, not adjacent to a metro area	0.0066 (0.0012)	1452	27345
(8)	Completely rural or less than 2,500 urban population, adjacent to a metro area	0.0019 (0.0015)	397	6549
(9)	Completely rural or less than 2,500 urban population, not adjacent to a metro area	0.0034 (0.0022)	351	5614

Notes: Each specification use a different sub-sample from the set of census tracts having information on house prices estimated by Bogin et al. (2016). Sub-samples are defined according to a measures of rurality at the county level ranging from 1 through 9 (see `https://dsdr-kb.psc.isr.umich.edu/answer/1102`). For each regression, the independent variable is the percentage of the area of a given census tract having potential wind capacity in 2008 greater than 30%, interacted with a dummy set to one in the years after the first wind turbine was installed in the US. All specifications include year fixed effects and standard errors are clustered at the census tract level.

Appendix Table 2: Monotonicity test, Zillow Home Value Index

|	△turbines	Coefficient (Standard error)	Number of locations	Number of location-time
(1)	Counties in metro areas of 1 million population or more	0.0002 (0.0001)	5322	1300018
(2)	Counties in metro areas of 250,000 to 1 million population	0.0003 (0.0001)	3234	773909
(3)	Counties in metro areas of fewer than 250,000 population	0.0004 (0.0001)	1870	440196
(4)	Urban population of 20,000 or more, adjacent to a metro area	0.0003 (0.0001)	1240	287556
(5)	Urban population of 20,000 or more, not adjacent to a metro area	0.0005 (0.0002)	295	61341
(6)	Urban population of 2,500 to 19,999, adjacent to a metro area	0.0010 (0.0003)	1098	234722
(7)	Urban population of 2,500 to 19,999, not adjacent to a metro area	0.0009 (0.0002)	448	86698
(8)	Completely rural or less than 2,500 urban population, adjacent to a metro area	0.0009 (0.0009)	160	31564
(9)	Completely rural or less than 2,500 urban population, not adjacent to a metro area	0.0007 (0.0005)	90	13829

Notes: Each specification use a different sub-sample from the set of ZIP codes having Zillow Home Value Index data. Sub-samples are defined according to a measures of rurality at the county level ranging from 1 through 9 (see `https://dsdr-kb.psc.isr.umich.edu/answer/1102`). For each regression, the independent variable is the percentage of the area of a given ZIP code having potential wind capacity in 2008 greater than 30%, interacted with a dummy set to one in the years after the first wind turbine was installed in the US. All specifications include calendar month fixed effects and standard errors are clustered at the ZIP code level.

Appendix Figure 1: Monthly house prices (ZHVI)

Notes: The sample is all ZIP code-months with data on the Zillow Home Value Index. The figure shows the cumulative distribution function for monthly log ZHVI truncated at $1 million.

Appendix Figure 2: Geographic coverage of BDL house price index

Notes: The sample is the set of ZIP codes having data on Bogin et al. (2016)'s house price index (HPI) in 2000 and 2005. Darker areas represent ZIP codes experiencing a larger increase in HPI between 2000 and 2005.

Appendix Figure 3: Estimates of potential wind capacity with different technology

Panel A: Current estimate vs 2008

slope = .847
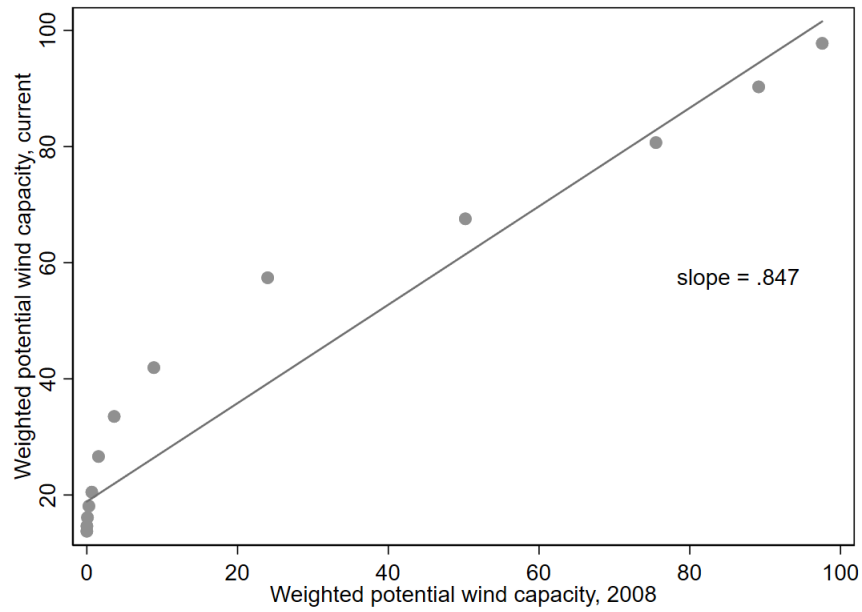
*Panel B: Near future estimate vs 2008*

slope = .717

Notes: Each panel reports the binned scatter plot between measures of wind capacity based on technology available in two distinct points in time. Potential wind capacity is estimated by Lopez et al. (2012). I divide ZIP codes into twenty equal-sized bins and plot, for each bin, the average potential wind capacity with the corresponding technology. The solid lines show a fit from an OLS regression on the underlying data. The slope of the estimated coefficients is reported in each plot.

Appendix Figure 4: Wind power class
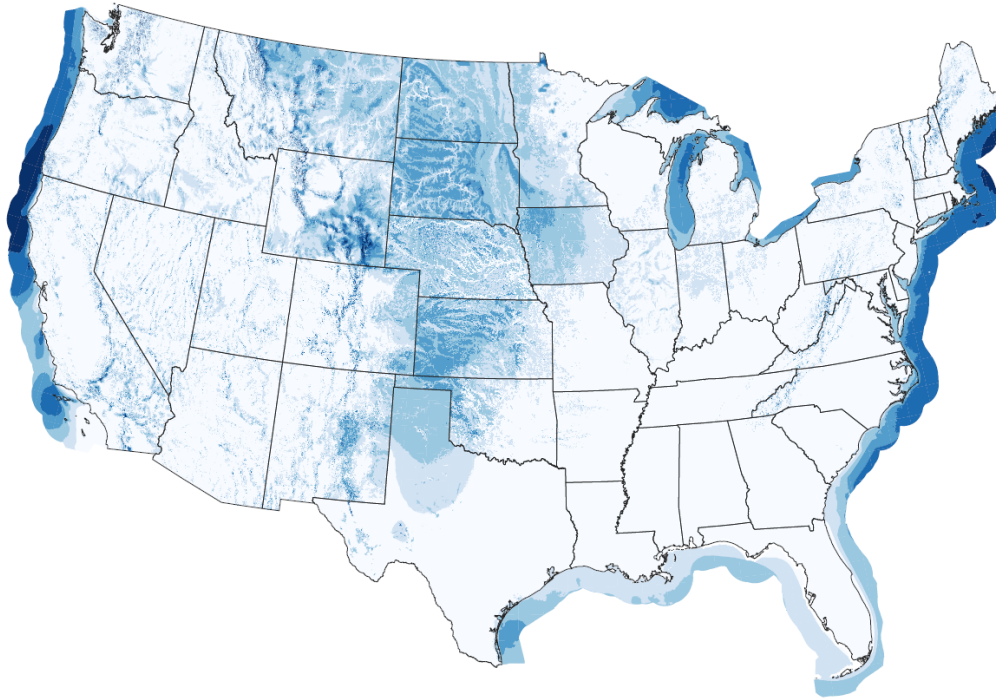
Notes: The figure reports the annual average wind power class (no exclusions applied) used in the Renewable Electricity Futures Study (`https://www.nrel.gov/analysis/re-futures.html`). Wind power class ranges from 1 (not shaded) to 7 (shaded in dark blue).

Appendix Figure 5: Number of currently active executives

Notes: Green triangles represent offshore wind turbines (not included in the main analysis)